

# Table of Contents

|                                          |          |
|------------------------------------------|----------|
| <b>Table of Contents</b>                 | <b>1</b> |
| <b>PFT Testing criteria</b>              | <b>2</b> |
| Goal                                     | 2        |
| Assumptions                              | 2        |
| Methodology                              | 2        |
| Requirements                             | 3        |
| Requirements for non HW offload datapath | 3        |
| Requirement for the HW offload datapath  | 4        |
| Pass table for 10Gbps interfaces         | 4        |
| Pass table for 25Gbps interfaces         | 5        |
| Pass table for 40Gbps interfaces         | 5        |
| Pass table for 50Gbps interfaces         | 6        |
| Pass table for 100Gbps interfaces        | 6        |
| References                               | 6        |

# PFT Testing criteria

## Goal

Define a technical requirement for the forwarding throughput measured on our products (RHEL and OVS+DPDK), to reduce the risk of customers claiming degradation in the forwarding performance when the same applications (typically NFVs) run on different PFs (i.e. NICs provided by different vendors).

## Assumptions

- different NIC models can negotiate different link speeds (10/25/40/50/100 Gb/s)
- when the traffic is limited by CPU resources on the host, the forwarding throughput is (much) lower than the theoretical frame rate that can be obtained at a given link speed.
- the maximum forwarding rate that can be obtained on the SW datapath (OVS+DPDK or OVS+Kernel SW) depends mostly on the driver implementation, and on computing resources in the host CPU.
- when HW offload is used, the maximum forwarding traffic is limited by computing resources on the NICs.

## Methodology

The PVP script [1] is used to benchmark the forwarding performance of a NIC model. The following configurations are tested:

| throughput               | component being tested               |
|--------------------------|--------------------------------------|
| -----+-----              |                                      |
| OVS+DPDK, L2 and L3      | poll-mode driver                     |
| OVS+Kernel SW, L2 and L3 | kernel driver                        |
| OVS+Kernel HW, L2 and L3 | kernel driver + hardware (TC Flower) |

With the aim of increasing the test coverage and underline potential limitations in the NIC hardware or in the NIC driver, tests are run using multiple frame sizes (64B to 1514B) and multiple simultaneous flows (10 to 1000000). Resulting values (in pps) are compared to the maximum theoretical rate for the given speed, or with values measured on a reference architecture under the same test conditions.

# Requirements

Each test is run on y-stream RHEL kernels, with latest FDB OVS+DPDK, following the instructions available at [1]

## Requirements for non HW offload (sw only path) datapath

### VSPerf DPDK

- 64 bytes PVP OVS-DPDK will achieve 3 Mpps at 0 loss for 10 minutes
- 1500 bytes PVP OVS-DPDK will achieve 1.5 Mpps at 0 loss for 10 minutes
- 64 bytes PVP OVS-DPDK multi-queue 2 Queue 4 PMD will achieve 6 Mpps at 0 loss for 10 minutes
- 1500 bytes PVP OVS-DPDK multi-queue 2 Queue 4 PMD will achieve 1.5 Mpps at 0 loss for 10 minutes
- 2000 byte jumbo frames PVP OVS-DPDK will achieve 1.1 Mpps at 0 loss for 10 minutes
- 9000 byte jumbo frames PVP OVS-DPDK will achieve 250 Kpps at 0 loss for 10 minutes

### VSPERF Kernel

- 64 bytes PVP OVS kernel will achieve 100 Kpps at 0.002 loss for 10 minutes
- 1500 bytes PVP OVS kernel will achieve 100 Kpps at 0.002 loss for 10 minutes

### VSPERF SR-IOV

- 64 Bytes PVP SR-IOV passthrough will achieve 10 Mpps at 0 loss for 10 minutes
- 1500 Bytes PVP SR-IOV passthrough will achieve 1.6 Mpps at 0 loss for 10 minutes

### PVP Tests

- Open vSwitch DPDK datapath 10 hour long duration (PASS if OVS is still passing traffic after 10 hours, without any crashes in between)
- OVS Linux Kernel datapath 10 hour long duration (PASS if OVS is still passing traffic after 10 hours, without any crashes in between)

### Beaker Functional Tests China OVS Team

- OVS kernel functional tests, needs to PASS all tests

Note that DPDK performance will be validated via VSperf and not via PVP.

## Requirement for the HW offload datapath

For ovs\_pvp it should pass the 10-hour long duration test, both L2, and L3. Which means it is still passing traffic after 10 hours, without any crash in between.

In addition, the 10 hours 10k L3 flows test should reach the minimum packets per second (pps) forwarding rate for each packet size specified to pass. The minimum pps rate per packet can be taken from the tables below. Each interface speed has its own table and is based on average numbers of already existing solutions.

Finally in addition to the above, the following flow characteristics should be met:

1. it must be possible to offload simultaneously up to 10K flows
2. the flow installation rate should not be less than 1.5K flows/s

| Packet Size | Minimum Packets Per Second |
|-------------|----------------------------|
| 64          | 2490513                    |
| 128         | 2400604                    |
| 256         | 2256274                    |
| 512         | 2114661                    |
| 768         | 1427664                    |
| 1024        | 1077586                    |
| 1514        | 733376                     |

**Pass table for 10Gbps interfaces**

| Packet Size | Minimum Packets Per Second |
|-------------|----------------------------|
| 64          | 9637204                    |
| 128         | 6508473                    |
| 256         | 4000375                    |
| 512         | 3693513                    |
| 768         | 3462503                    |
| 1024        | 2671544                    |
| 1514        | 1819386                    |

**Pass table for 25Gbps interfaces**

| Packet Size | Minimum Packets Per Second |
|-------------|----------------------------|
| 64          | 10217438                   |
| 128         | 9092933                    |
| 256         | 8657281                    |
| 512         | 6536093                    |
| 768         | 5280149                    |
| 1024        | 4024205                    |
| 1514        | 2619575                    |

**Pass table for 40Gbps interfaces**

| Packet Size | Minimum Packets Per Second |
|-------------|----------------------------|
| 64          | 26067795                   |
| 128         | 23473241                   |
| 256         | 18772297                   |
| 512         | 10476825                   |
| 768         | 7083181                    |
| 1024        | 5347950                    |
| 1514        | 3640232                    |

**Pass table for 50Gbps interfaces**

| Packet Size | Minimum Packets Per Second |
|-------------|----------------------------|
| 64          | 26899051                   |
| 128         | 24033668                   |
| 256         | 18772297                   |
| 512         | 13542611                   |
| 768         | 9978782                    |
| 1024        | 7948678                    |
| 1514        | 5451593                    |

**Pass table for 100Gbps interfaces**

## References

[1] PVP setup, [https://github.com/ctrautma/RHEL\\_NIC\\_QUALIFICATION](https://github.com/ctrautma/RHEL_NIC_QUALIFICATION)